

docker

Introduction to Containers

By: Nitish Jadia

Outline

- Container recipe
- Why should I care?
- A quick docker demo
- Building blocks
- Security

Outline

- **Container recipe**
- Why should I care?
- A quick docker demo
- Building blocks
- Security

Containers? These?



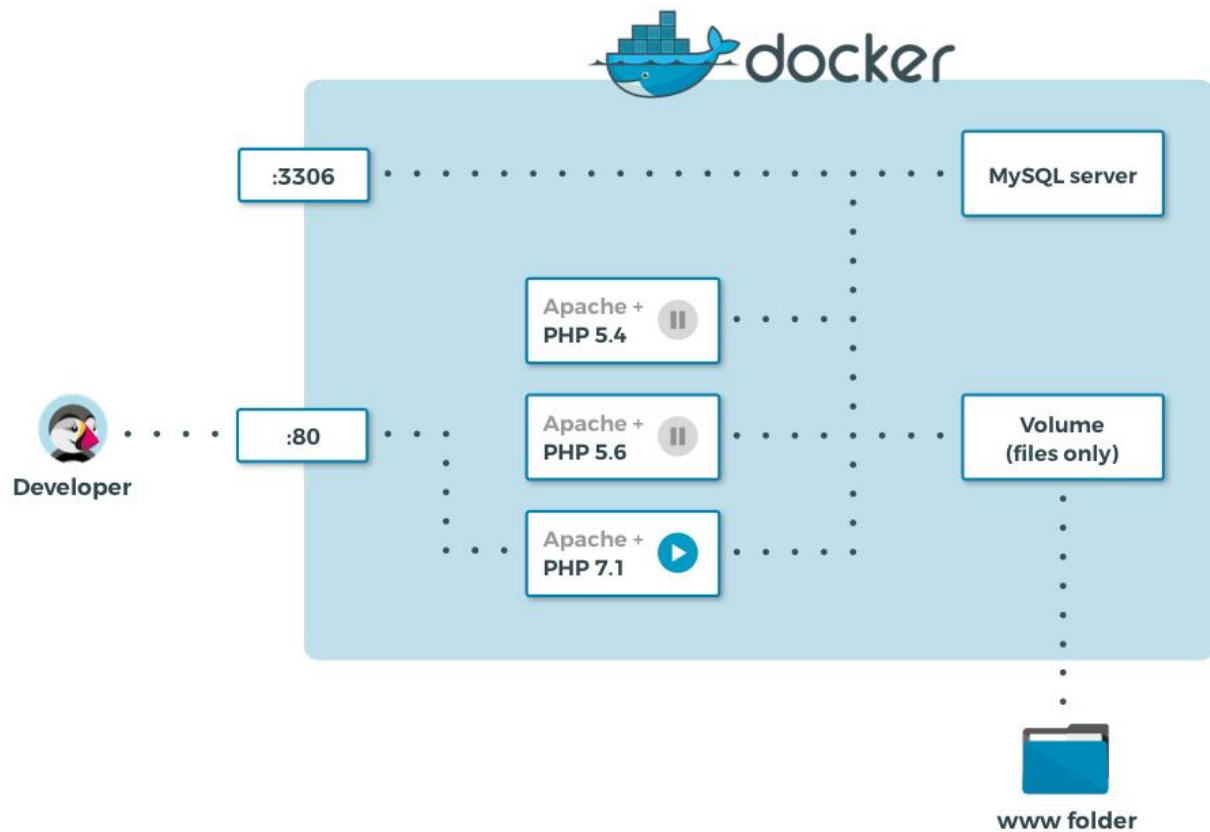
Before containers ...things were messy



B 4433



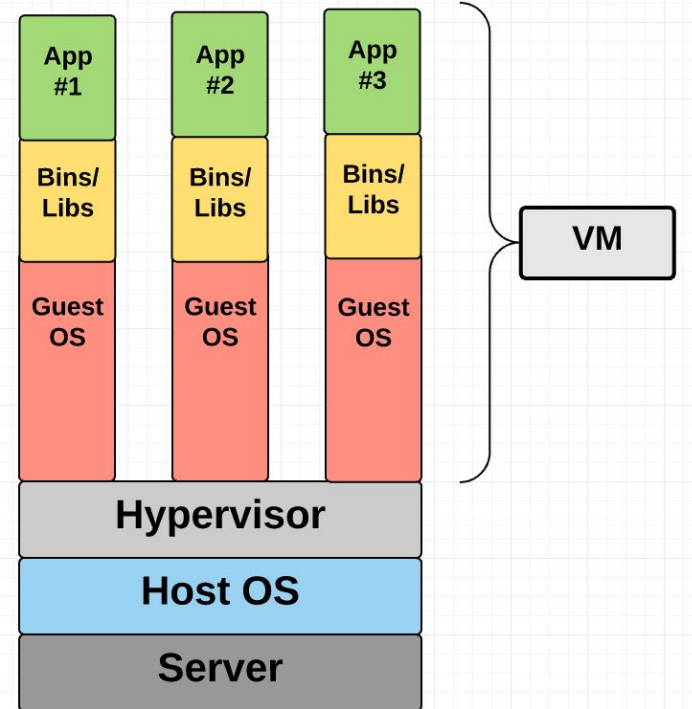
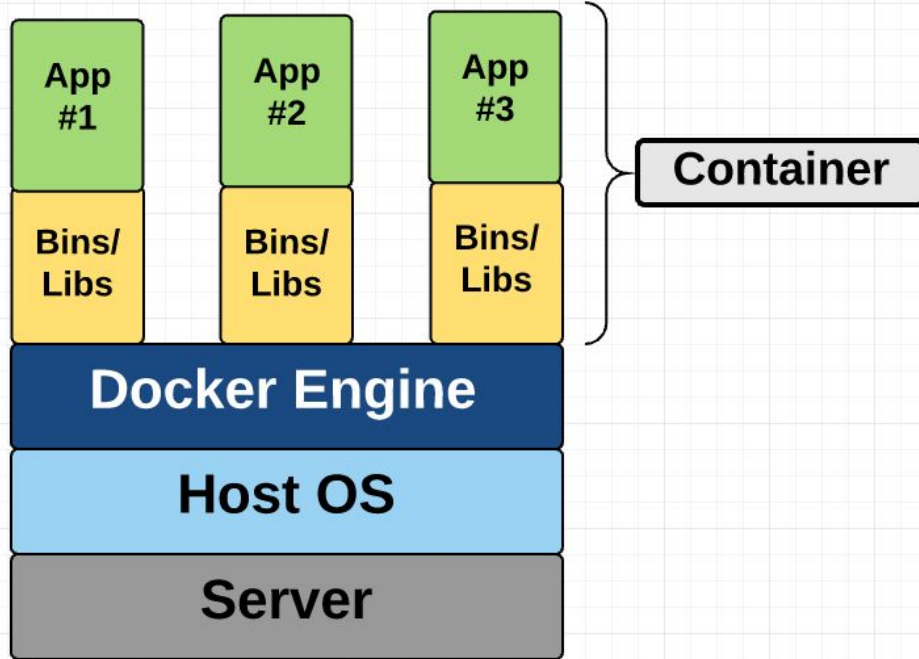
source: <https://en.wikipedia.org/>



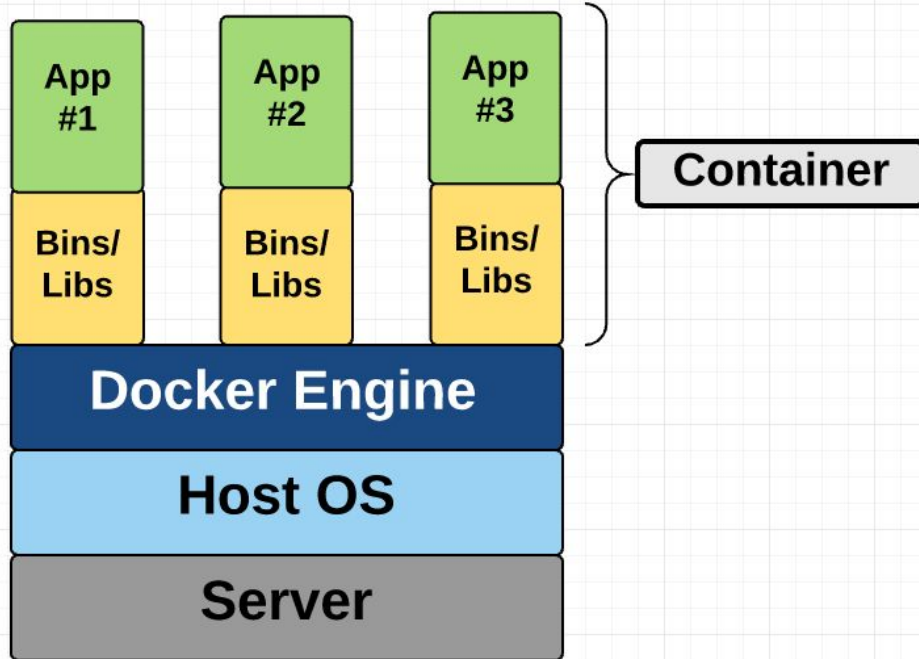
source: <http://build.prestashop.com/assets/images/2017/02/prestashop-docker.jpg>

Organise your server with containers

Containers



Containers





source: <https://medium.com>

- All containers share the same kernel of the host system.
Pro: Extremely reduced performance overhead.
- Better utilization of resources due to shared kernel.
- Lightweight and uses less space on disk.
- Portable and better dependency management

Outline

- Container recipe
- **Why should I care?**
- A quick docker demo
- Building blocks
- Security

Deploy anywhere and anything



- webapps
- backends
- SQL, NoSQL
- big data
- load balancing
- ... and more



... but, it was working on my machine.

Deploy reliably & consistently

- If it works locally, it will work on the server
- With exactly the same behavior
- Regardless of versions
- Regardless of distros
- Regardless of dependencies
- Typical laptop runs 10-100 containers easily
- Typical server can run 100-1000 containers

Outline

- Container recipe
- Why should I care?
- **A quick docker demo**
- Building blocks
- Security

Demo

- Docker installation
- Docker CLi
 - Basic commands
 - Pull images
 - Deploy containers
- Docker hub

Is docker running?

```
docker run hello-world
```



File Edit View Search Terminal Help
rover@metal1:~\$ docker run hello-world

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/get-started/>

docker run

```
docker run -ti debian bash
```

-ti -> terminal interactive



rover@metal1: ~

File Edit View Search Terminal Help

```
rover@metal1:~$ docker run -ti debian bash
```

```
Unable to find image 'debian:latest' locally
```

```
latest: Pulling from library/debian
```

```
741437d97401: Pull complete
```

```
Digest: sha256:066051f6674f6a3293bbd5a190081b1ae7fcae655a3884db59ebb3a2831da623
```

```
Status: Downloaded newer image for debian:latest
```

```
root@07681df639c2:/#
```

Check list of images

docker images



```
rover@metal1:~$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
<none>	<none>	12e7d41f679b	2 weeks ago	411MB
rover_wp	latest	378620ce8cc4	2 weeks ago	411MB
wordpress	php7.1-apache	378620ce8cc4	2 weeks ago	411MB
mysql	5.7	e47e309f72c8	2 weeks ago	372MB
ubuntu	latest	47b19964fb50	2 weeks ago	88.1MB
debian	latest	d508d16c64cd	2 weeks ago	101MB
hello-world	latest	fce289e99eb9	7 weeks ago	1.84kB

```
rover@metal1:~$
```

Pulling images

```
docker pull ubuntu:xenial
```



```
rover@metal1: ~  
File Edit View Search Terminal Help  
rover@metal1:~$ docker pull ubuntu:xenial  
xenial: Pulling from library/ubuntu  
7b722c1070cd: Pull complete  
5fbf74db61f1: Pull complete  
ed41cb72e5c9: Pull complete  
7ea47a67709e: Pull complete  
Digest: sha256:e4a134999bea4abb4a27bc437e6118fdddfb172e1b9d683129b74d254af51675  
Status: Downloaded newer image for ubuntu:xenial  
rover@metal1:~$
```


Terminology

Images - The blueprints of our application which form the basis of containers. In the demo above, we used the `docker pull` command to download the busybox or ubuntu image.

Containers - Created from Docker images and run the actual application. We create a container using `docker run` which we did using the busybox image that we downloaded. A list of running containers can be seen using the `docker ps` command.

Docker Daemon - The background service running on the host that manages building, running and distributing Docker containers. The daemon is the process that runs in the operating system to which clients talk to.

Docker Client - The command line tool that allows the user to interact with the daemon. More generally, there can be other forms of clients too - such as Kitematic which provide a GUI to the users.

Docker Hub - A registry of Docker images. You can think of the registry as a directory of all available Docker images. If required, one can host their own Docker registries and can use them for pulling images.

Run docker container from images

```
docker run -ti ubuntu:latest bash
```

ubuntu -> image

latest -> tag (optional, by default it's latest)

bash -> what do we want to do with the image.



```
root@4ca326454d31: /  
File Edit View Search Terminal Help  
rover@metal1:~$ docker run -ti ubuntu:latest bash  
root@4ca326454d31:/#
```

Leave container running in background(detatch)

```
docker run -d -ti ubuntu /bin/bash
```

-d -> detaches the container



```
rover@metall1:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

```
rover@metall1:~$ docker run -d -ti ubuntu /bin/bash
```

```
9bb79dace90e42d5d78d1df03e467536af33abcf9f36770ca1bf551bbb587e20
```

```
rover@metall1:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
9bb79dace90e	ubuntu	"/bin/bash"	7 seconds ago	Up 5 seconds		ecstatic_northcu

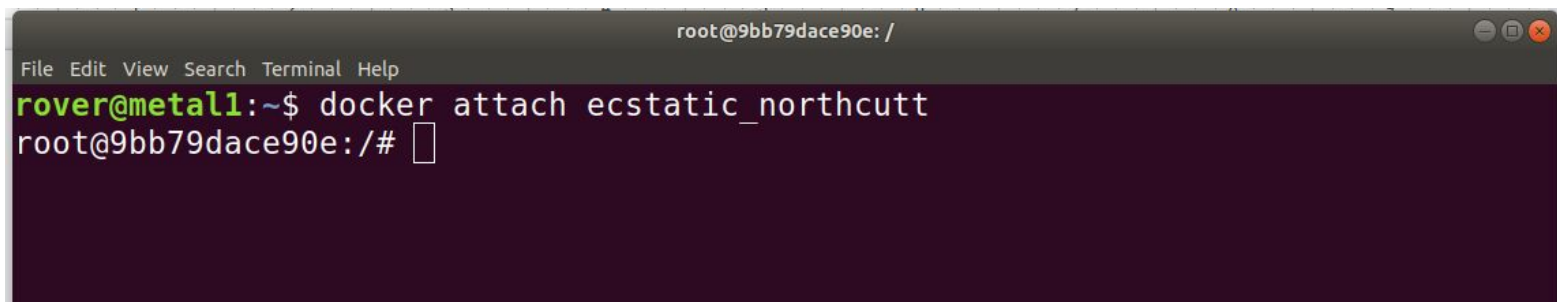
```
tt
```

```
rover@metall1:~$
```

Attach to a container

```
docker attach ecstatic_northcutt
```



A terminal window with a dark gray title bar and a dark purple background. The title bar contains the text 'root@9bb79dace90e: /' and three window control icons (minimize, maximize, close). The terminal has a menu bar with 'File', 'Edit', 'View', 'Search', 'Terminal', and 'Help'. The prompt 'rover@metal1:~\$' is in green and blue, followed by the command 'docker attach ecstatic_northcutt'. The next line shows the prompt 'root@9bb79dace90e:/#' with a white cursor box.

```
root@9bb79dace90e: /
File Edit View Search Terminal Help
rover@metal1:~$ docker attach ecstatic_northcutt
root@9bb79dace90e:/#
```


Additional commands

- `docker ps -a`
- `docker info`
- `docker restart zealous_darwin`
- `docker inspect blisful_saha`
- `docker inspect blisful_saha | grep -i ip`



Dockerfile

- A Dockerfile is a simple text-file that contains a list of commands that the Docker client calls while creating an image.
- It's a simple way to automate the image creation process.

Let's create a Dockerfile:

```
mkdir build
```

```
cd build
```

```
vim Dockerfile
```

Dockerfile

```
#This is a custom ubuntu image with vim already installed
```

```
FROM ubuntu:xenial
```

```
MAINTAINER nitish <hello@nitishjadia.com>
```

```
RUN apt-get update
```

```
RUN apt-get install -y vim
```

Dockerfile

```
docker build -t="nitishmod/ubuntuvim:v3" .
```

-t -> title

. -> dot, because Dockerfile is in the same folder.

```
rover@metal1:~$ mkdir build
rover@metal1:~$ cd build/
rover@metal1:~/build$ vim Dockerfile
rover@metal1:~/build$ #Dockerfile content pasted
rover@metal1:~/build$ docker build -t="nitishmod/ubuntuvim:v3" .
Sending build context to Docker daemon 2.048kB
Step 1/4 : FROM ubuntu:xenial
--> 7e87e2b3bf7a
Step 2/4 : MAINTAINER nitish <hello@nitishjadia.com>
--> Running in 9f0ad6983572
Removing intermediate container 9f0ad6983572
--> 05064e105607
Step 3/4 : RUN apt-get update
--> Running in 8ec408dfe347
Get:1 http://security.ubuntu.com/ubuntu xenial-security InRelease [109 kB]
Get:2 http://security.ubuntu.com/ubuntu xenial-security/main amd64 Packages [786 kB]
Get:3 http://security.ubuntu.com/ubuntu xenial-security/restricted amd64 Packages [12.7 kB]
Get:4 http://security.ubuntu.com/ubuntu xenial-security/universe amd64 Packages [541 kB]
Get:5 http://archive.ubuntu.com/ubuntu xenial InRelease [247 kB]
Get:6 http://security.ubuntu.com/ubuntu xenial-security/multiverse amd64 Packages [6116 B]
Get:7 http://archive.ubuntu.com/ubuntu xenial-updates InRelease [109 kB]
Get:8 http://archive.ubuntu.com/ubuntu xenial-backports InRelease [107 kB]
```

File Edit View Search Terminal Help

```
Setting up file (1:5.25-2ubuntu1.1) ...
Setting up libexpat1:amd64 (2.1.0-7ubuntu0.16.04.3) ...
Setting up libmpdec2:amd64 (2.4.2-1) ...
Setting up libssl1.0.0:amd64 (1.0.2g-1ubuntu4.14) ...
debconf: unable to initialize frontend: Dialog
debconf: (TERM is not set, so the dialog frontend is not usable.)
debconf: falling back to frontend: Readline
debconf: unable to initialize frontend: Readline
debconf: (Can't locate Term/ReadLine.pm in @INC (you may need to install the Term::ReadLine module) (@INC contains: /etc/perl /usr/local
/lib/x86_64-linux-gnu/perl/5.22.1 /usr/local/share/perl/5.22.1 /usr/lib/x86_64-linux-gnu/perl5/5.22 /usr/share/perl5 /usr/lib/x86_64-lin
ux-gnu/perl/5.22 /usr/share/perl/5.22 /usr/local/lib/site_perl /usr/lib/x86_64-linux-gnu/perl-base .) at /usr/share/perl5/Debconf/FrontE
nd/Readline.pm line 7.)
debconf: falling back to frontend: Teletype
Setting up libpython3.5-minimal:amd64 (3.5.2-2ubuntu0~16.04.5) ...
Setting up mime-support (3.59ubuntu1) ...
Setting up libsqlite3-0:amd64 (3.11.0-1ubuntu1) ...
Setting up libpython3.5-stdlib:amd64 (3.5.2-2ubuntu0~16.04.5) ...
Setting up vim-common (2:7.4.1689-3ubuntu1.2) ...
Setting up libpython3.5:amd64 (3.5.2-2ubuntu0~16.04.5) ...
Setting up vim-runtime (2:7.4.1689-3ubuntu1.2) ...
Setting up vim (2:7.4.1689-3ubuntu1.2) ...
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/vim (vim) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/vimdiff (vimdiff) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/rvim (rvim) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/rview (rview) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/vi (vi) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/view (view) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/ex (ex) in auto mode
update-alternatives: using /usr/bin/vim.basic to provide /usr/bin/editor (editor) in auto mode
Processing triggers for libc-bin (2.23-0ubuntu10) ...
Removing intermediate container dc86af9e3fc3
--> 5430a6e937b1
Successfully built 5430a6e937b1
Successfully tagged nitishmod/ubuntuvim:v3
rover@metal1:~/build$
```

```
rover@metal1:~/build$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
nitishmod/ubuntuvim	v3	5430a6e937b1	2 minutes ago	200MB
<none>	<none>	12e7d41f679b	2 weeks ago	411MB
rover_wp	latest	378620ce8cc4	2 weeks ago	411MB
wordpress	php7.1-apache	378620ce8cc4	2 weeks ago	411MB
mysql	5.7	e47e309f72c8	2 weeks ago	372MB
ubuntu	latest	47b19964fb50	2 weeks ago	88.1MB
debian	latest	d508d16c64cd	2 weeks ago	101MB
ubuntu	xenial	7e87e2b3bf7a	4 weeks ago	117MB
hello-world	latest	fce289e99eb9	7 weeks ago	1.84kB

```
rover@metal1:~/build$
```


Apache web server

```
> mkdir /home/rover/public_html
```

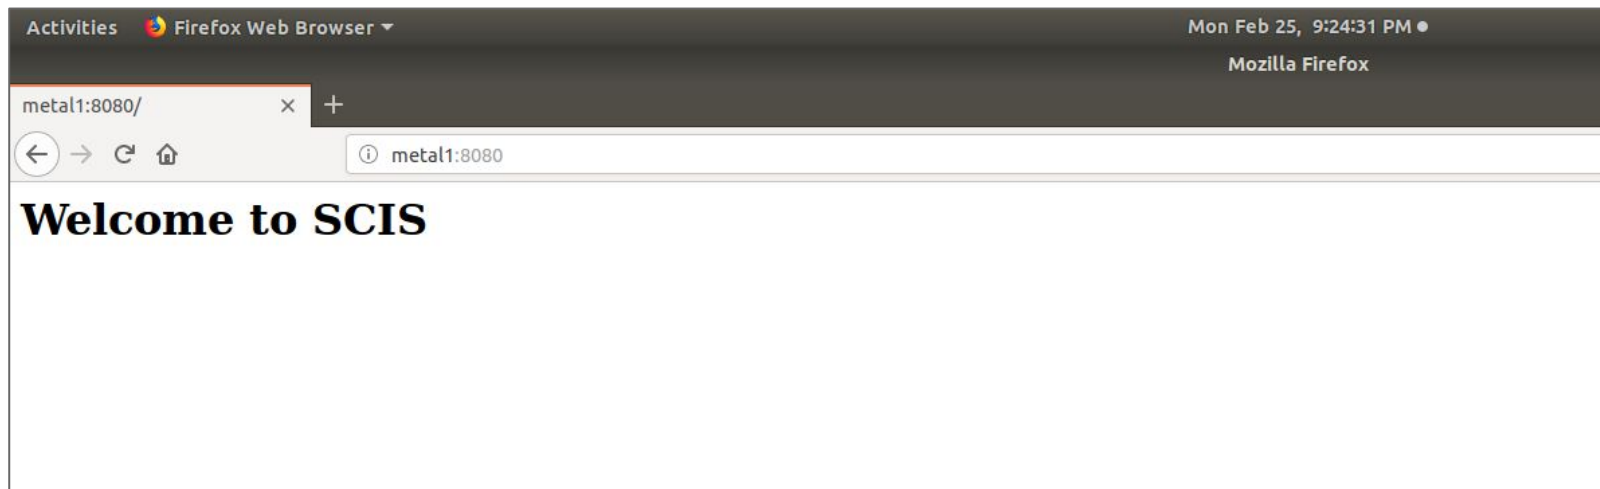
```
> sudo docker run -dit --name webserver-test -p 8080:80 -v  
/home/user/website:/usr/local/apache2/htdocs/ httpd:2.4
```

```
> cd /home/rover/public_html
```

```
> echo ""  
<html> <head>  
<h1> Welcome to SCIS </h1>  
</head> </html> "" > index.html
```

```
Activities Terminal
Mon Feb 25, 9:33:53 PM
rover@metal1: ~/public_html

rover@metal1:~$ mkdir public_html
rover@metal1:~$ sudo docker run -dit --name webserver-test -p 8080:80 -v /home/rover/public_html:/usr/local/apache2/htdocs/ httpd:2.4
Unable to find image 'httpd:2.4' locally
2.4: Pulling from library/httpd
6ae821421a7d: Already exists
0ceda4df88c8: Pull complete
24f08eb4db68: Pull complete
ddf4fc318081: Pull complete
fc5812428ac0: Pull complete
Digest: sha256:5e7992fcdad214d5e88c4dfde274befe60d5d5b232717862856012bf5ce31086
Status: Downloaded newer image for httpd:2.4
c8df1dc8e72f59c1db09b2e0379ed94e3bba59c7692eeebca9380d8fa58d4813
rover@metal1:~$ cd /home/rover/public_html/
rover@metal1:~/public_html$ echo ""
> <html>
> <head>
> <h1> Welcome to SCIS </h1>
> </head> </html> "" > index.html
rover@metal1:~/public_html$
```



Docker Hub

- Docker Hub is a cloud-based repository in which Docker users and partners create, test, store and distribute container images.
- <https://hub.docker.com/>
- Users can host and share their own custom images
- Apache HTTP server image: https://hub.docker.com/_/httpd



httpd ☆

Docker Official Images

The Apache HTTP Server Project

10M+

Container Linux x86-64 IBM Z PowerPC 64 LE 386 ARM 64 ARM Application Infrastructure

Official Image

Linux - x86-64 (latest)

Copy and paste to pull this image

```
docker pull httpd
```



[View Available Tags](#)

DESCRIPTION

REVIEWS

TAGS

Supported tags and respective Dockerfile links

- 2.4.38 , 2.4 , 2 , latest (2.4/Dockerfile)
- 2.4.38-alpine , 2.4-alpine , 2-alpine , alpine (2.4/alpine/Dockerfile)

Quick reference

- Where to get help:**
the Docker Community Forums, the Docker Community Slack, or Stack Overflow
- Where to file issues:**
<https://github.com/docker-library/httpd/issues>



Please log in to write a review of this product.

Outline

- Container recipe
- Why should I care?
- A quick docker demo
- **Building blocks**
- Security

Building blocks of containers

- Again.. what is a container?
- Cgroups
- Namespaces

Building blocks of containers

- **Again.. what is a container?**
- Cgroups
- Namespaces

What is a container?

- How it "feels" like:
 - own process space
 - own network space
 - run stuff as root
 - can install packages
 - can run services
 - can mess up routing, iptables ...

What is a container?

- It's not quite like a VM:
 - uses the host kernel
 - can't boot a different OS
 - can't have its own modules
 - doesn't need init as PID 1
 - doesn't need syslogd, cron...
- It's just normal processes on the host machine
 - contrast with VMs which are opaque

Building blocks of containers

- Again.. what is a container?
- **Cgroups**
- Namespaces



Control Groups (Cgroups)

- Resource metering and limiting
 - memory
 - CPU
 - block I/O
 - network*
- Device node (/dev/*) access control
- Crowd control

Control Groups (Cgroups)

- Each subsystem has a hierarchy (tree)
 - separate hierarchies for CPU, memory, block I/O...
- Hierarchies are independent
 - the trees for e.g. memory and CPU can be different
- Each process is in a node in each hierarchy
 - think of each hierarchy as a different dimension
- Each hierarchy starts with 1 node (the root)
- Each node = group of processes
 - sharing the same resources

Control Groups (Cgroups)



Outline

- Container recipe
- Why should I care?
- A quick docker demo
- **Building blocks**
- Security

Building blocks of containers

- Again.. what is a container?
- Cgroups
- Namespaces

Building blocks of containers

- **Again.. what is a container?**
- Cgroups
- Namespaces

What is a container?

- How it "feels" like:
 - own process space
 - own network space
 - run stuff as root
 - can install packages
 - can run services
 - can mess up routing, iptables ...

What is a container?

- It's not quite like a VM:
 - uses the host kernel
 - can't boot a different OS
 - can't have its own modules
 - doesn't need init as PID 1
 - doesn't need syslogd, cron...
- It's just normal processes on the host machine
 - contrast with VMs which are opaque

Building blocks of containers

- Again.. what is a container?
- **Cgroups**
- Namespaces



Control Groups (Cgroups)

- Resource metering and limiting
 - memory
 - CPU
 - block I/O
 - network*
- Device node (/dev/*) access control
- Crowd control

Control Groups (Cgroups)

- Each subsystem has a hierarchy (tree)
 - separate hierarchies for CPU, memory, block I/O...
- Hierarchies are independent
 - the trees for e.g. memory and CPU can be different
- Each process is in a node in each hierarchy
 - think of each hierarchy as a different dimension
- Each hierarchy starts with 1 node (the root)
- Each node = group of processes
 - sharing the same resources

Control Groups (Cgroups)



Memory cgroup: accounting

- Keeps track of pages used by each group:
 - file (read/ write from block devices)
 - active (recently accessed)
 - inactive (candidate for eviction)
- Each page is "charged" to a group
- Page can be shared across multiple groups
 - e.g. multiple processes reading from the same files when pages are shared, only one group "pays" for a page.

Memory cgroup: accounting

- Keeps track of pages used by each group:
 - file (read/ write from block devices)
 - active (recently accessed)
 - inactive (candidate for eviction)
- Each page is "charged" to a group
- Page can be shared across multiple groups
 - e.g. multiple processes reading from the same files when pages are shared, only one group "pays" for a page.

Memory cgroup: limits

- Each group can have its own limits
 - limits are optional
 - two kinds of limits:
 - Soft limits
 - Hard limits
- Soft limits are not enforced
 - they influence reclaim memory pressure

Memory cgroup: limits

- Hard limits will trigger a per-group OOM killer
- Limits can be set for different kinds of memory
 - physical memory
 - kernel memory
 - total memory
- Multiple groups use the same page, only first one is "charged"
 - but if it stops using it, the charge is moved to another group.

CPU cgroup

- Keep track of user/system CPU time
- Keeps track of usage per CPU
- Allows to set weights

CPUset cgroup

- Pin groups to specific CPU(s)
- Reserve CPUs for specific apps
- Avoid processes bouncing between CPUs

Blkio cgroup

- Keeps track of I/O for each group
 - per block device
 - read vs write
- Set throttle (limits) for each group
 - per block device
 - read vs write

net_cls and net_prio cgroup

- Automatically set traffic class or priority, for traffic generated by processes in the group
- Only works for egress traffic
- Net_cls will assign traffic to a class
 - class then has to be matched with tc/iptables, otherwise traffic just flows normally
- Net_prio will assign traffic to a priority
 - priorities are used by queuing disciplines

Devices cgroup

- Controls what the group can do on devices nodes
- Permissions include read/write
- Typical use:
 - allow `/dev/{tty,zero,random,null}`
 - deny everything else
- A few interesting nodes that we can expose:
 - `/dev/net/tun` (network interface manipulation)
 - `/dev/fuse` (filesystems in use space)
 - `/dev/dri` (GPU)

Freezer cgroup

- Allows to freeze a group of processes
- Similar in functionality to SIGSTOP
- Cannot be detected by the processes
 - that's how it differs from SIGSTOP
- Can be used in process migration

Building blocks of containers

- Again.. what is a container?
- Cgroups
- **Namespaces**

Namespaces

- Provide processes with their own system view
 - Cgroups = limits how much you can use;
 - Namespaces = limits what you can see
- Multiple namespaces:
 - pid
 - net
 - mnt
 - uts
 - ipc
 - user
- Each process is in one namespace of each type

pid namespace

- Processes within a PID namespace only see processes in the same PID namespace
- Each PID namespace has its own numbering
 - starting at 1
- If PID 1 goes away, whole namespace is killed
- Those namespaces can be nested
- A process ends up having multiple PIDs
 - one per namespace in which its nested

Network namespaces: in theory

- Processes within a given network namespace get their own private network stack, including:
 - network interfaces
 - routing tables
 - iptable rules
 - sockets

Network namespaces: in practice

- Use virtual interfaces acting as a cross-over cable
- eth0 in container network namespace, paired with vethXX in host network namespace.
- All the vethXX are bridged together via virtual switch inside the container host.
 - Docker calls the bridge docker0

Mnt namespace

- Process can have their own root fs
 - conceptually close to chroot
- Processes can also have "private" mounts
 - /tmp
 - masking of /proc, /sys
- Mounts can be totally private, or shared

Uts namespace

- Let containers set their own hostname.
 - `gethostname/ sethostname`

IPC namespace

- Allows a process (or group of processes) to have own:
 - IPC semaphores
 - IPC message queues
 - IPC shared memory
- without risk of conflict with other instances.

User namespace

- Allows to map UID.
 - UID 0 in container --> UID 9999 in host
 - UID 1000 in container --> UID 1564 in host
- UID in containers becomes irrelevant
- Security improvement

Outline

- Container recipe
- Why should I care?
- A quick docker demo
- Building blocks
- **Security**

Use RunC Flaw to gain Root access on Host

- The vulnerability, identified as **CVE-2019-5736**, was discovered by two open source security researchers on 11th Feb 2019.
- “High level” container runtimes like Docker does image creation and management
- Use "Low level" runC to handle tasks related to running containers
 - creating a container
 - attaching a process to an existing container (docker exec)

The Vulnerability

Overview given by the runC team:

The vulnerability allows a malicious container to (with minimal user interaction) overwrite the host runc binary and thus gain root-level code execution on the host. The level of user interaction is being able to run any command ... as root within a container in either of these contexts:

- Creating a new container using an attacker-controlled image.
- Attaching (docker exec) into an existing container which the attacker had previous write access to.

Bibliography

1. What is container? | <https://www.docker.com/resources/what-container>
2. Demystifying containers 101 | <https://tinyurl.com/yxd7fnpe>
3. Getting started with Docker | <https://tinyurl.com/y35e879j> , <http://bit.do/eJwmy>
4. Docker for development and deployment | <http://bit.do/eJwmH>
5. RunC Flaw Lets Attackers Escape Linux Containers to Gain Root on Hosts | <https://is.gd/3qPVsQ>
6. CVE-2019-5736 | <https://nvd.nist.gov/vuln/detail/CVE-2019-5736>
7. Explaining CVE-2019-5736 | <http://bit.do/eJLcG>
8. LXC | <https://linuxcontainers.org/lxc/introduction/> , <https://is.gd/i0F9i4>
9. Container basics | <https://is.gd/Eq4gaP> <https://is.gd/pGgMqW>
10. Namespaces in operation | <https://lwn.net/Articles/531114/>
11. Future of Linux Containers | <https://is.gd/U0QR9K>
12. cgroups | <https://www.kernel.org/doc/Documentation/cgroup-v1/cgroups.txt> , <https://is.gd/1x1ihc>
13. Attacks on Linux Containers | <https://ieeexplore.ieee.org/document/7854163>

Bibliography

14. Comparing different containers | <https://ieeexplore.ieee.org/document/8075934>
15. Beginners Guide to Containers Technology and How it Actually Works[video] | <https://is.gd/CzTb1d>
16. My Docker Guide | <https://github.com/w4rb0y/devOpsNotes/blob/master/dockerGuide.md>
17. Google Cloud | <https://cloud.google.com/containers/>
18. Containers from scratch | <https://ericchiang.github.io/post/containers-from-scratch/>

Thank You

Query?

Container without Docker

- Setup a file system
- chroot
- unshare
- nsenter
- bind mounts
- cgroups

chroot

A **chroot** is an operation that changes the apparent root directory for the current running process and their children. A program that is run in such a modified environment cannot access files and commands outside that environmental directory tree. This modified environment is called a **chroot jail**.

chroot

Download the rootfs

wget <https://is.gd/nEYUgv>

sudo tar -xzf rootfs.tar.gz

sudo chroot rootfs /bin/bash

mount -t proc proc /proc

pkill top

chroot

```
ldd /bin/bash | awk '{print $3}' | sed '/^$/d'
```

cgroups

```
ls /sys/fs/cgroup/  
sudo su
```

```
mkdir /sys/fs/cgroup/memory/demo
```

```
ls /sys/fs/cgroup/memory/demo/
```

unshare

```
sudo chroot rootfs /bin/bash
```

```
sudo unshare -p -f --mount-proc=$PWD/rootfs/proc \
```

```
chroot rootfs /bin/bash
```

-p -> pid

-f -> Fork the specified program as a child process of unshare rather than running it directly. This is useful when creating a new PID namespace.

Thank You